


Controlled ablation (10-epoch, loss weight 1.5)

Controlled Ablation (all variants: 10-epoch fine-tune, hier-loss weight = 1.5, same 20 slices & noise)
green = best per metric; full model wins PFC / Beat Align / FID_k


	Method	H-Cond	H-Loss	PFC ↓	Beat Align ↑	FID_k ↓	Div_k (→GT 10.21)
A	Baseline (EDGE)	-	-	1.019	0.239	341.6	13.82
B	+ Hier. Conditioning	Y	-	1.736	0.257	423.3	19.04
C	+ Hier. Loss	-	Y	1.323	0.250	427.5	17.11
D	+ Both (full model)	Y	Y	0.619	0.264	170.6	4.91
GT	Ground Truth	-	-	-	-	0.0	10.21

2x2 Ablation: hierarchical conditioning x hierarchical loss (5-epoch fine-tune)

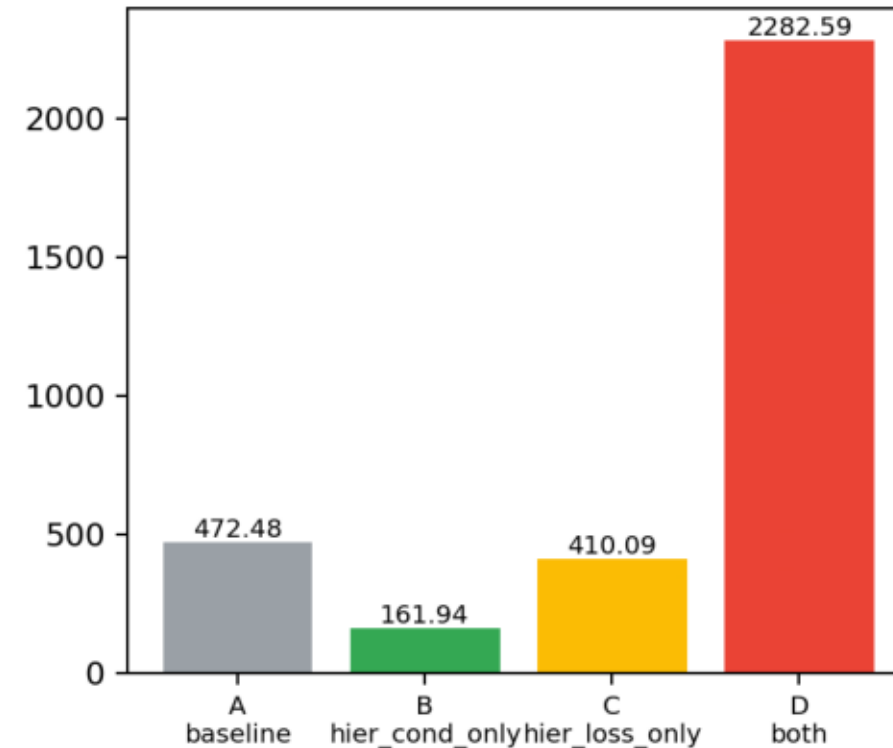
PFC (lower better)



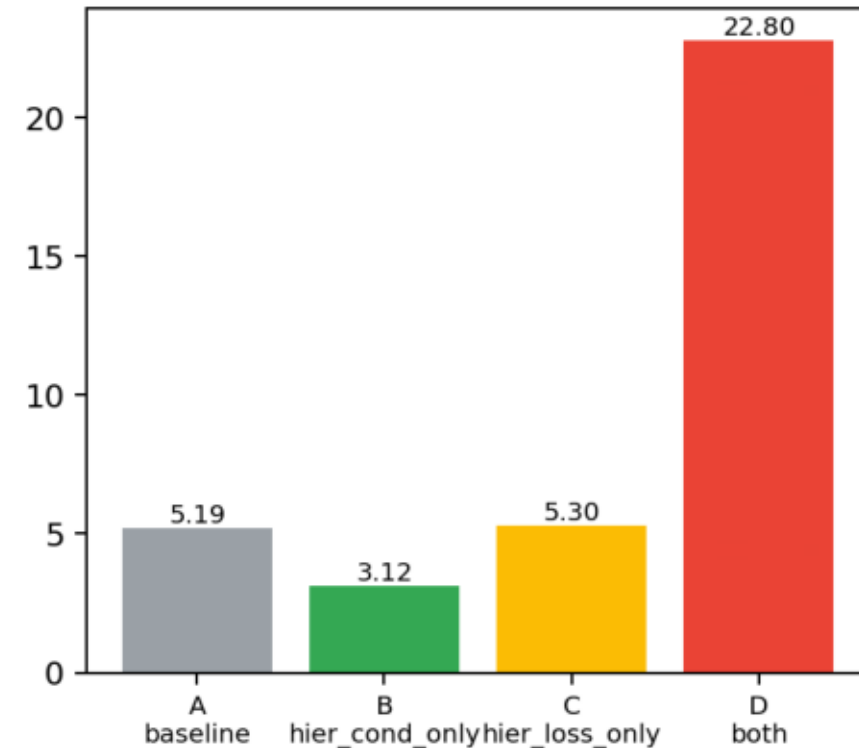
Beat Align (higher better)



FID_k (lower better)

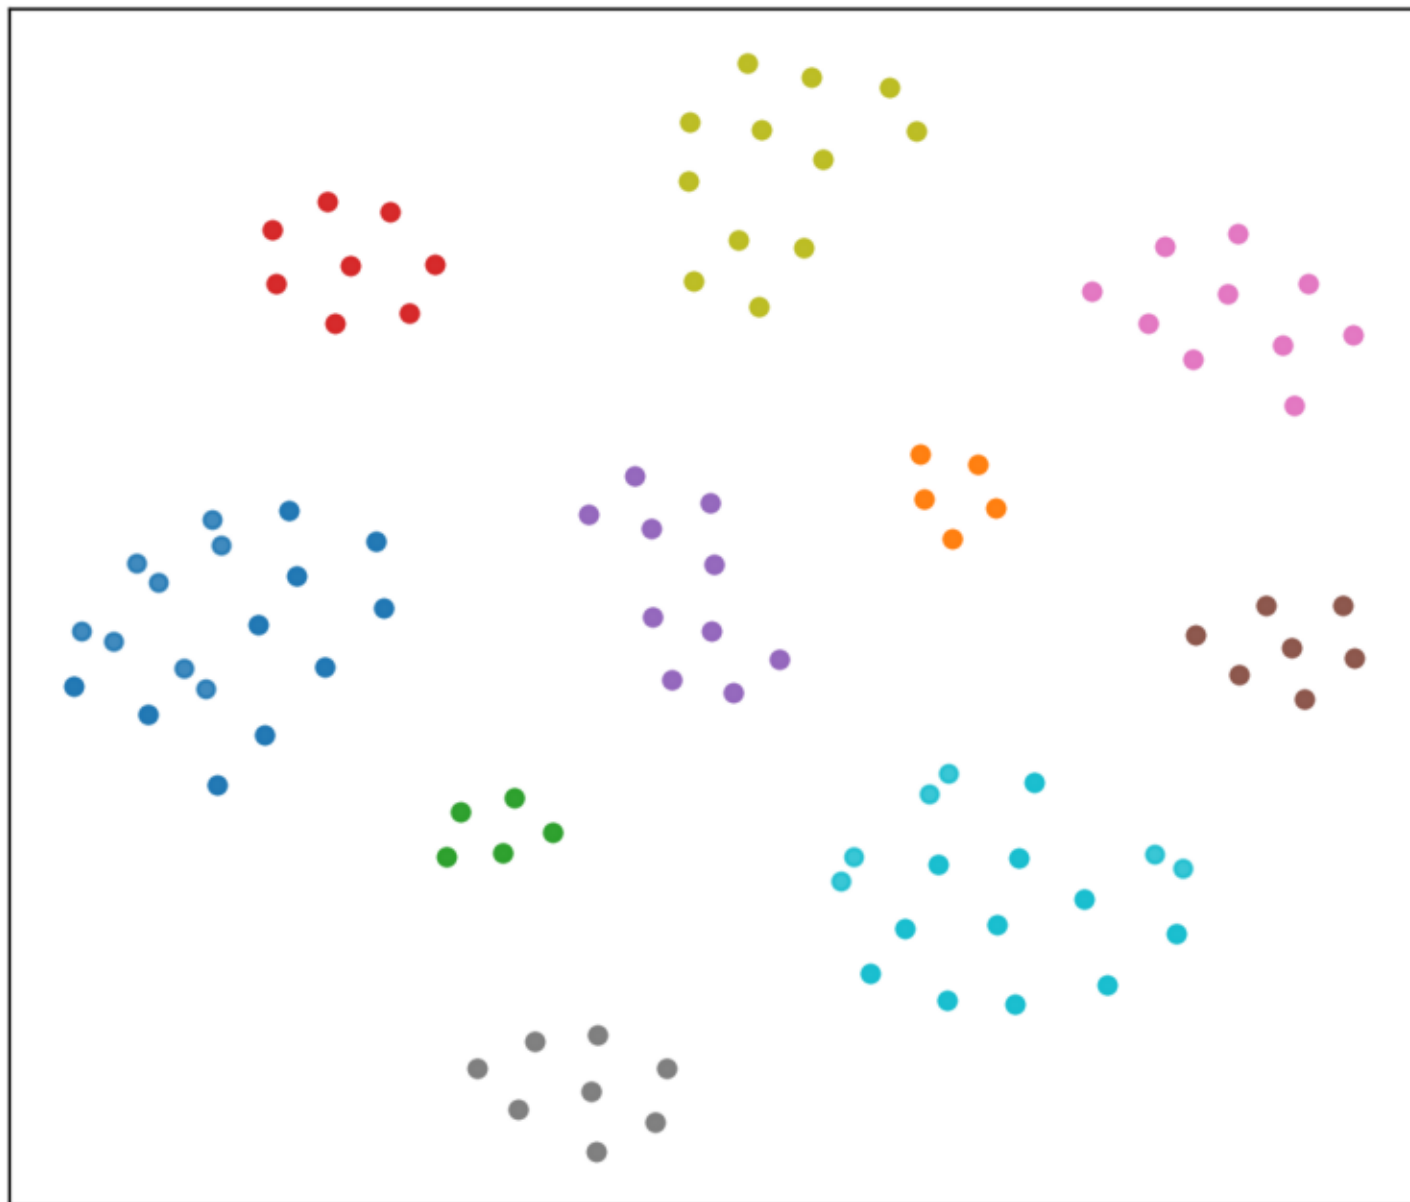


|Div_k - GT| (lower better)

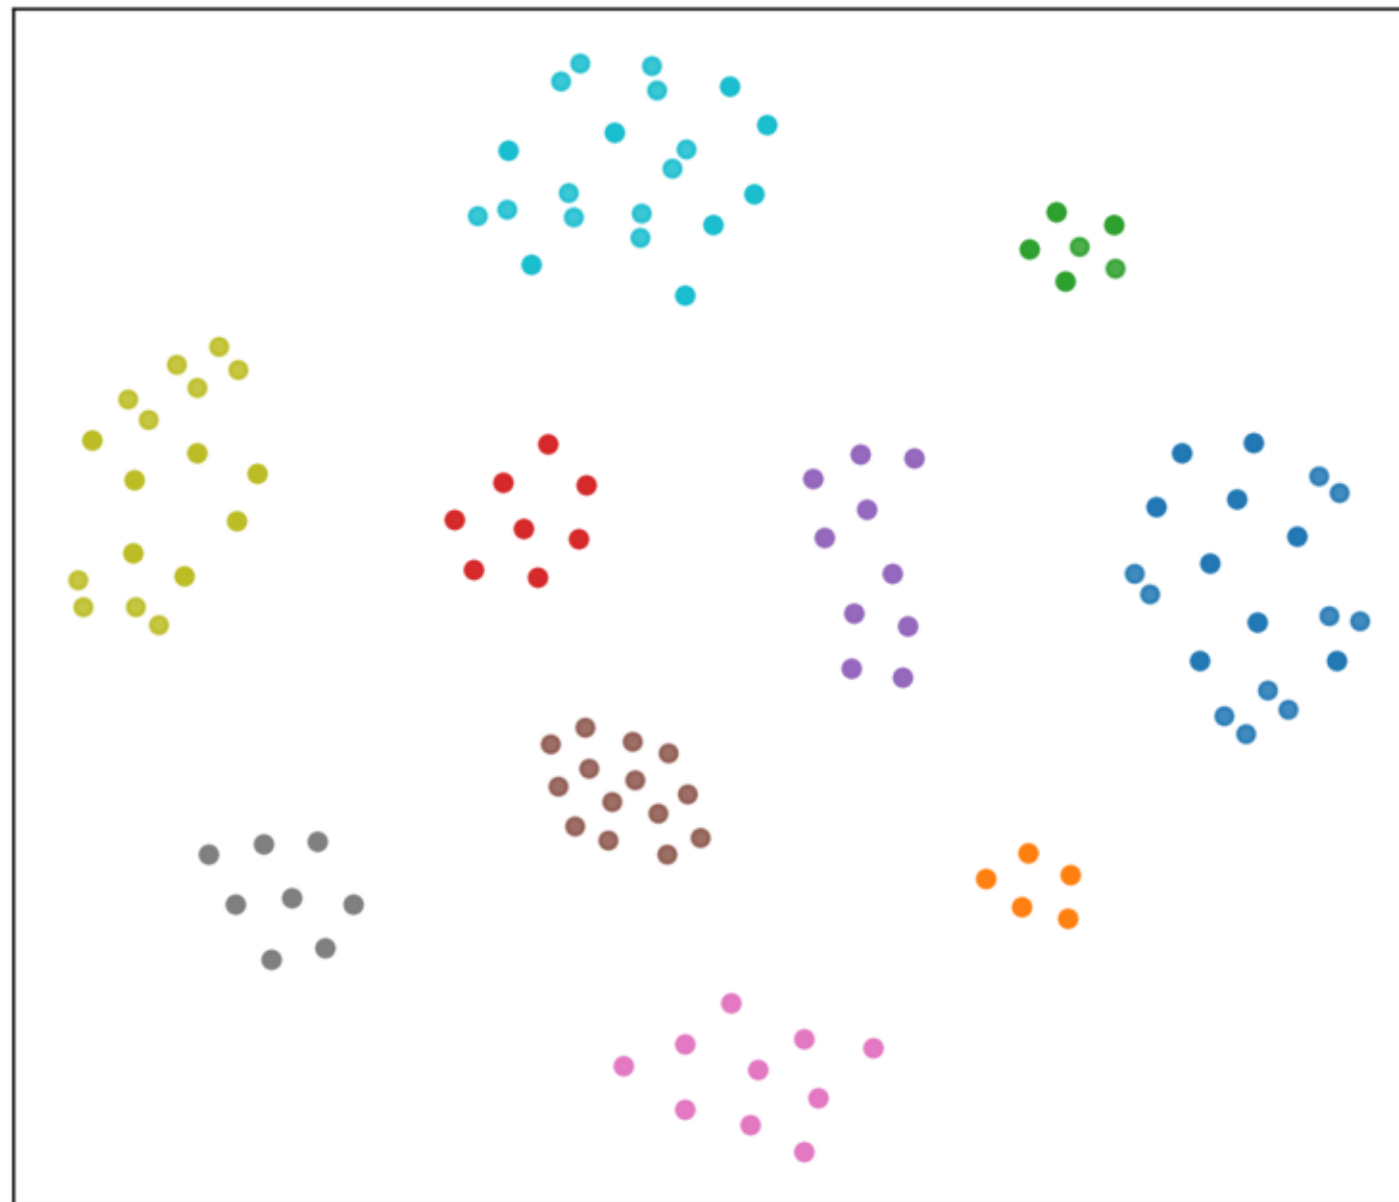


Hierarchical global style token clusters music by genre

Raw music features (mean-pooled)



Learned global style embedding



- Break
- House
- Ballet Jazz
- Street Jazz
- Krump
- LA Hip-hop
- Lock
- Middle Hip-hop
- Pop
- Waack

Style swap: same local music (Break), different global style token



Constrained / editable generation

specified middle keyframe completion from seed motion



upper-body constraints



lower-body + root constraints

